

**TO  
THE  
NEW™**



**Terraform**

Akshay Verma

- => Terraform - What and Why?
- => A Quick view.
- => Launching EC2 with comparing against Cloudformation.
- => State Files: Glance
- => S3 backend.
- => Use of Dynamodb.
- => Remote State (How data is fetched)
- => terraform workspace Vs Infra Environment
- => Terraform Functions
- => Modularization and Variabilization. Why and How?
- => Directory Structure Planning, a few cases.
- => Questions and Cases



**KEEP  
CALM  
AND  
AUTOMATE ALL  
THE THINGS**

## What:

=> Terraform is a tool for building, changing, and versioning infrastructure safely and efficiently.

## Why:

=> Terraform is cool.

=> Infrastructure as a Code.

=> It is completely platform agnostic.

=> A single tool to manage virtual clouds(AWS, VMWare, Azure), supporting services like DNS, Email, or managing some administration in your database.

=> Terraform's speed and operations are exceptional. Plan actually allows us to see what's gonna change.

# What is Terraform

---



Terraform is a tool for building, changing, and versioning infrastructure safely and efficiently. Terraform can manage existing and popular service providers as well as custom in-house solutions.

Configuration files describe to Terraform the components needed to run a single application or your entire datacenter. Terraform generates an execution plan describing what it will do to reach the desired state, and then executes it to build the described infrastructure. As the configuration changes, Terraform is able to determine what changed and create incremental execution plans which can be applied.

<https://www.terraform.io/intro/index.html>

<https://learn.hashicorp.com/terraform?track=getting-started#getting-started>

<https://www.terraform.io/docs/providers/aws/>

<https://github.com/terraform-aws-modules?>

# Installation



## 1. Install unzip

```
sudo apt-get install unzip
```

## 2. Download latest version of the terraform

```
wget https://releases.hashicorp.com/terraform/0.12.24/terraform_0.12.24_linux_amd64.zip
```

## 3. Extract the downloaded file archive

```
unzip terraform_0.12.24_linux_amd64.zip
```

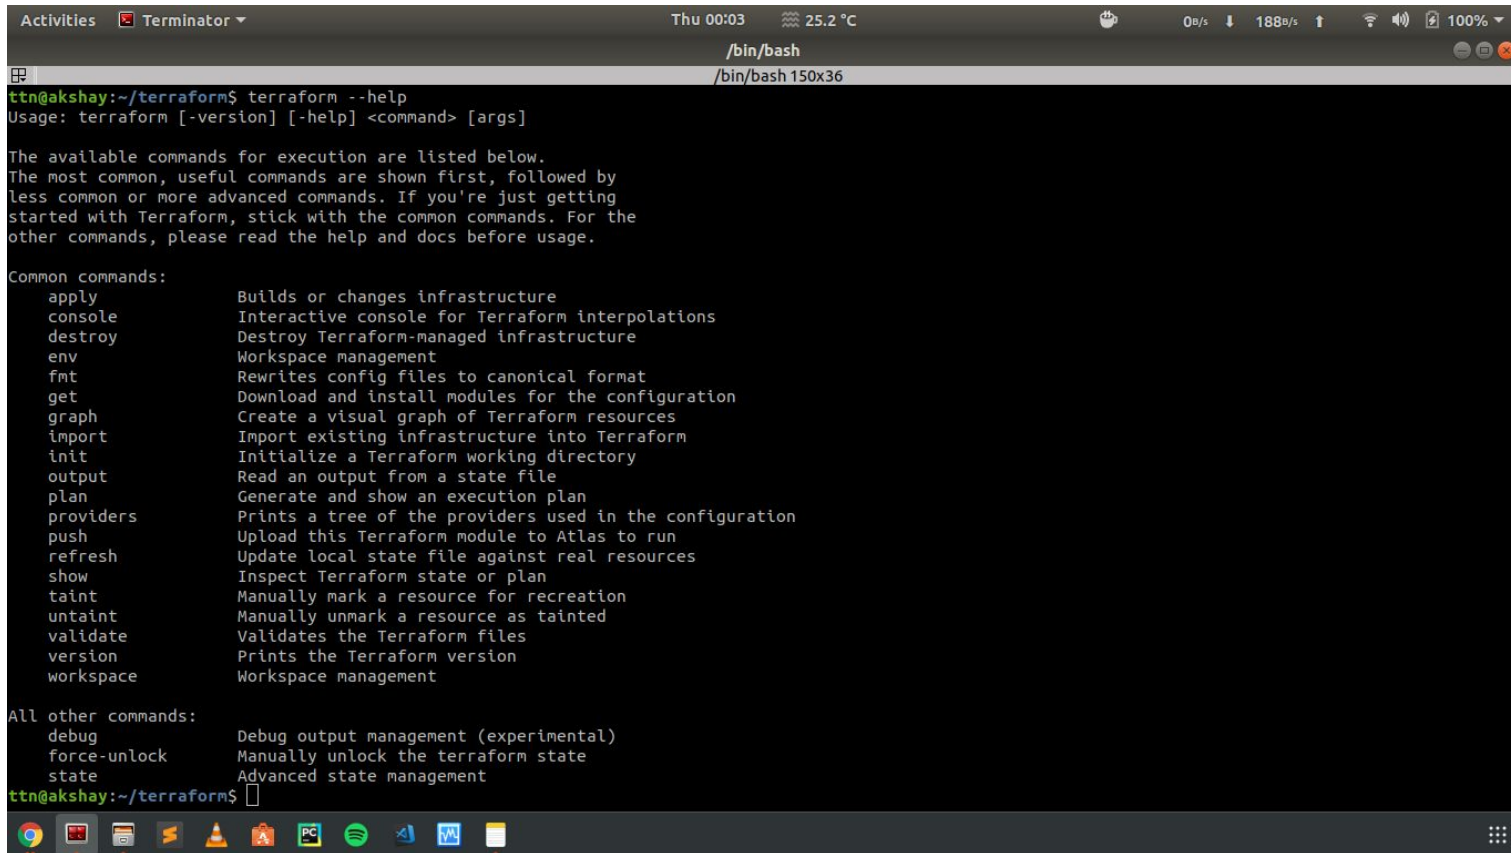
## 4. Move the executable into a directory searched for executables

```
sudo mv terraform /usr/local/bin/
```

## 5. Run it

```
terraform --version
```

# Terraform Commands



```
Activities Terminator Thu 00:03 25.2 °C 0s/s 188s/s 100%  
/bin/bash  
ttn@akshay:~/terraform$ terraform --help  
Usage: terraform [-version] [-help] <command> [args]  
  
The available commands for execution are listed below.  
The most common, useful commands are shown first, followed by  
less common or more advanced commands. If you're just getting  
started with Terraform, stick with the common commands. For the  
other commands, please read the help and docs before usage.  
  
Common commands:  
  apply          Builds or changes infrastructure  
  console        Interactive console for Terraform interpolations  
  destroy        Destroy Terraform-managed infrastructure  
  env            Workspace management  
  fmt            Rewrites config files to canonical format  
  get            Download and install modules for the configuration  
  graph          Create a visual graph of Terraform resources  
  import         Import existing infrastructure into Terraform  
  init           Initialize a Terraform working directory  
  output         Read an output from a state file  
  plan           Generate and show an execution plan  
  providers      Prints a tree of the providers used in the configuration  
  push           Upload this Terraform module to Atlas to run  
  refresh        Update local state file against real resources  
  show           Inspect Terraform state or plan  
  taint          Manually mark a resource for recreation  
  untaint        Manually unmark a resource as tainted  
  validate       Validates the Terraform files  
  version        Prints the Terraform version  
  workspace      Workspace management  
  
All other commands:  
  debug          Debug output management (experimental)  
  force-unlock   Manually unlock the terraform state  
  state          Advanced state management  
ttn@akshay:~/terraform$
```

# How Does Terraform look like?

```
provider "aws" {  
    access_key = "ACCESS_KEY_HERE"  
    secret_key = "SECRET_KEY_HERE"  
    region     = "us-east-1"  
}
```

The name to reference within Terraform files.

```
resource "aws_instance" "test" {  
    ami           = "ami-07ebfd5b3428b6f4d"  
    instance_type = "t2.micro"
```

```
    tags = {  
        Name = "terraform-test"  
        owner = "akshay.verma@tothenew.com"  
        purpose = "bootcamp-tf"  
    }  
}
```

The actual name of bucket and other properties.

Here, the combination "aws\_instance" and "test" must be unique.

The provider block is used to configure the named provider, in our case "aws". A provider is responsible for creating and managing resources. Multiple provider blocks can exist if a Terraform configuration is composed of multiple providers, which is a common situation.

The resource block defines a resource that exists within the infrastructure. A resource might be a physical component such as an EC2 instance



## Terraform

```
resource "aws_instance" "Nginx_Server" {  
  ami                = "${var.ami_id}"  
  instance_type      = "${var.instance_type}"  
  subnet_id          = "${var.subnet_id}"  
  vpc_security_group_ids = ["${var.security_groups}"]  
  key_name            = "${var.key_name}"  
  iam_instance_profile = "${var.role_name}"  
  ebs_optimized       = true  
  tags                = "${var.tags}"  
  volume_tags         = "${var.tags}"  
}
```

## Cloudformation

```
"Nginx_Server":{  
  "Type":"AWS::EC2::Instance",  
  "Properties":{  
    "ImageId": {"Ref":"ESAMIId"},  
    "InstanceType": {"Ref":"ESInstanceType"},  
    "IamInstanceProfile" : { "Ref": "EC2ProfileForTag" },  
    "KeyName": { "Ref":"KeyName"},  
    "Tags":[  
      {"Key":"Name", Value": { "Fn::Join":["-",[{"Ref": "Envv"},  
        "nginx-server"]}]}}  
    ]  
  }  
}
```

## Variables to override the defaults

Variable.tf


```
variable "env" {}
```

# So how you do it

---

First You See the Plan..

Then You Apply..



# Creating Infrastructure

---



- terraform init
- terraform plan -var env=\$env -out=tfplan
- terraform apply "tfplan"
- terraform destroy -var env=\$env

# Lets Dive Terraform

# State and State Files



=> It is a JSON file which looks like:

# State File

```
{  
  "version": 3,  
  "terraform_version": "0.11.13",  
  "serial": 112,  
  "lineage": "dcb1eebd-13c6-929e-4928-0486ffa9adb4",  
  "modules": [  
    { "resources": {  
      "aws_instance.utility": {  
        "type": "aws_instance",  
        "depends_on": [],  
        "primary": {  
          "id": "i-0ce13100a7a2a2ed2",  
          "attributes": {  
            "ami": "ami-0cfe7d2728d424846",  
            "arn": "arn:aws:ec2:ap-south-1:349554833603:instance/i-0ce13100a7a2a2ed2",  
            "associate_public_ip_address": "false",  
            "availability_zone": "ap-south-1b" ..... (contd)          }  
        }  
      }  
    }  
  ]  
}
```

# State and State Files

---

- => Terraform stores metadata and is needed for syncing to real world infrastructure.
- => Terraform must store state about your managed infrastructure and configuration.
- => This state is stored by default in a local file named "terraform.tfstate"
- => Terraform uses this local state to create plans and make changes to your infrastructure.
- => Prior to any operation, Terraform does a refresh to update the state with the real infrastructure.
- => It is a JSON file which looks like:



# S3 backend and State Locking

## Remote State

- => Allows multiple members of a team work on the same Terraform code.
- => Latest State File is always available to each member.
- => Remote Datastore: s3, consul etc

```
terraform {  
  backend "s3" {  
    bucket = "my_project_terraformbackend"  
    key = "my_project/terraform.tfstate"  
    region = "ap-south-1"  
  }  
}
```

# Why This Key?



```
key = "my_project/terraform.tfstate"
```

## State Locking:

- => No two members should run TF at same time.
- => Prevents broken state file usage.
- => Uses DynamoDB to store lock state at any point of time.
- => Create a table with key type string named "LockID" which is also a hash key.

```
terraform {  
  backend "s3" {  
    bucket = "my_project_terraformbackend"  
    key = "my_project/terraform.tfstate"  
    region = "ap-south-1"  
    dynamodb_table = "terraform-lock"  
  }  
}
```

# Fetching data via Remote State



Create Resource in one terraform and use it everywhere??

Like import and export

YES its possible.. Let's see how.

# Fetching data via Remote State

1. Export the resource of source TF that needs to be Used. ==>

## Source Terraform

```
output "alb-https-listner-arn" {  
  value = "${module.app-lb.https-listner-arn}"  
}
```

2. Define the remote state details in destination terraform ==>

The details are of source TF.

## Destination Terraform

```
data "terraform_remote_state" "common" {  
  backend = "s3"  
  workspace = "<Workspace of source TF directory"  
  config {  
    bucket = "tatasky-wallet-terraform-tfstate"  
    key = "<Key of source TF directory"  
    region = "ap-south-1"  
  }  
}
```

3. Extract and Use the exported value as ==>

```
${data.terraform_remote_state.common.alb-https-listner-arn}
```

# Workspace and Infrastructure



- => Need to launch exact similar resources without affecting the current one, workspace is your savior.
- => As with other applications, workspace is a segregation between different versions of the same configuration.
- => Can compare with python environments
- => Different States can be maintained without coinciding with one another.
- => Works with a set of "terraform workspace" commands.  
<https://www.terraform.io/docs/state/workspaces.html>

## How this can be useful?

Workspace Vs Infra Environments

Workspace Vs AWS accounts

Workspace Vs Projects

# Terraform Functions

---



To transform and combine values and to do a whole lot of different things.

**Numeric Functions:** `abs()`, `ceil`, `floor`, `log`, `max`, `min` etc

**String:** `lower`, `upper`, `split`, `strrev`, `title`, `join`, `indent` etc

**Collection:** `concat`, `distinct`, `contains`, `element`, `index`, `keys`, `length`, `list`, `lookup`, `map`,

**Encoding:** `base64encode`, `base64decode`

**File System:** `File`, `path`

**Hash:** `md5`

**Conversion:** `tobool`, `tolist`, `tostring`

# Example



```
{for s in var.list : s => upper(s)}  
${file("${path.module}/config/filename.tpl")}  
map( "Name" , "${var.project_name_prefix}" )  
"${merge( var.tags, map( "Name" , "${var.project_name_prefix}" ) )}"  
Count = "${length(var.public_subnets)}"  
Cidr = "${lookup(var.public_subnets[count.index], "cidr")}"  
"${100 + (count.index * 100)}"  
  
vpc_id   = "${var.create_new_vpc == "true" ? join("",aws_vpc.vpc.*.id) :  
var.vpc_id}"
```

```
public_subnets =  
[  
  {  
    az = "ap-south-1a"  
    cidr = "10.23.34.0/27"  
  },  
  {  
    az = "ap-south-1b"  
    cidr = "10.23.34.32/27"  
  }  
]
```



# Terraform Modules



Modules are the key factors for reusable terraform code.

But What is a module?

Any set of Terraform configuration files in a folder is a module.

We can pass the values while calling a module to provide custom data.

**We all are using modules but are we really Modularizing?**

# Can this be treated as a module?

```
resource "aws_instance" "Nginx_Server" {  
  ami                = "${var.ami_id}"  
  instance_type      = "${var.instance_type}"  
  subnet_id          = "${var.subnet_id}"  
  vpc_security_group_ids = ["${var.security_groups}"]  
  key_name            = "${var.key_name}"  
  iam_instance_profile = "${var.role_name}"  
  ebs_optimized       = true  
  tags                = "${var.tags}"  
  volume_tags         = "${var.tags}"  
}
```

# A few examples of a good module



A good module must be a complete unit in itself. For example:

VPC: Creates complete VPC including dynamic subnets, NAT's, route tables etc.

ECS Service: Task Definition, Target Group, Service, Scaling and Alarms.

ASG: ASG, Scaling, Alarms, Lifecycle Hooks, Spot configurations etc.

Networking: CF, Route53, Lambda, WAF etc.

Should not use hard coded values.

Must contain all properties in a configurable form, i.e should allow all configurations to be passed externally.

Should not be dependent over any other external module.

# Variabilization

---



Important part to reuse the code.

Terraform generally use .tfvars files to provide variable values.

Keep everything possible to be variablize so that same code can be used just by changing the variable values.

# What structure are we using



- .modules
- terraform

```
|__ module_name (sub-dir name)  
  |__ aws.tf  
  |__ Main.tf / (module-wise segregation)  
  |__ outputs.tf  
  |__ variable.tf  
|__ <env.tfvars>
```

```
cd module_name;  
Terraform workspace select <env>  
Terraform apply -var-file=../<env.tfvars>
```

# Questions ?